ENHANCING STOCK PRICE PREDICTION WITH DEEP REINFORCEMENT LEARNING: A TIME SERIES ANALYSIS APPROACH

[1]AMEEN UR REHMAN, [2]TAHIR MAJID, [3]ADEEL JAVED, [4]MOHD SAIF AHMED, [5]MOHD ASHAD NAUSHAD, [6]MD ZUBAIR ALAM

[1,2,3,4,5,6]Department of Computer Science and Engineering, Jamia Hamdard, New Delhi, India
E-mail: ameensadiya12345@gmail.com, [2]tahir19.fm@gmail.com, [3]adeel.javed3511@gmail.com, [4]saifahmed9308@gmail.com, [5]mohda170602@gmail.com, [6]mdzubairalam.hamdard@gmail.com

Abstract - In the dynamic domain of financial markets, achieving precise forecasts of stock trends stands as an ongoing challenge. This research paper takes a profound dive into the captivating universe of Deep Reinforcement Learning (DRL) and its transformative role in reshaping stock prediction methodologies, intricately woven into the fabric of time series analysis. The paper involves thorough analysis and readings to explore the patterns and give relatively accurate results with a combined approach that uses traditional Time series techniques with Deep Reinforcement Learning. The DRL model used in the analysis is extremely effective in dealing with the complex data set of stock values and can overall enhance the accuracy of predictions as it is effective in sequential decision making it extremely useful in the constantly changing stock market. The use of DRL significantly reduces the overall dependency on manual feature engineering that is relatively time and resource-consuming by using deep neural networks that learn relevant data directly from the raw time data series.Additionally, the DRL model can effectively adapt to the ever-changing market without many constant manual adjustments thus decreasing the cost and overall maintenance of the feature, and therefore making it more reliable and effective than its counterparts. The results and insights from this research offer valuable contributions to the field of financial forecasting and open avenues for future research to explore even more sophisticated architectures and techniques in pursuit of ever more accurate predictions in dynamic stock markets.

Keywords - Data Science; Deep Reinforcement Learning; Machine Learning; Time Series; Deep Learning, Reinforcement Learning

I. INTRODUCTION

Predicting the behavior of highly volatile financial time series, such as stock returns at an intermediate frequency, is often challenging due to significant noise. However, enhancing forecasting accuracy is achievable by accurately predicting volatility or noise levels. This can be accomplished by integrating external variables and employing input features that encompass both market-specific information and broader contextual factors. The accuracy of these forecasts depends on various factors, including feature selection, observation frequency, normalization techniques, and the chosen model structure. While traditional deep learning (DL) approaches excel in processing complex input sets, they may not be suitable for scenarios requiring sequential decision-making [14][15]. Deep learning technology and Artificial Neural Networks (ANN) have gained widespread applications across diverse domains, such as healthcare, visual recognition, text analysis, cybersecurity, and more. They have significantly advanced these fields, although creating effective DL models for real-world, dynamic, and variable problems remains a challenge [15][16][17][18]. Moreover, traditional deep learning techniques tend to become black-box models, lacking fundamental understanding, which can impede progress at a foundational level. On the other hand, Reinforcement Learning (RL), while valuable for sequential decision-making, may not be optimal for complex datasets. RL focuses on algorithms making decisions to maximize cumulative rewards in specific environments, without relying on labeled input/output data. It emphasizes maintaining a balance between exploration and exploitation, with Q-learning being an off-policy reinforcement learning technique aiming to find the best actions based on the current state and future rewards. In the context of stock forecasting, this methodology can be applied to predict market trends. Deep Reinforcement Learning (DRL) bridges the strengths of both Reinforcement Learning (RL) and deep learning (DL) methodologies, making it more suitable for scenarios requiring sequential decision-making on complex input sets [16][18][19][20]. DRL emerges as a promising alternative for predicting stock returns, addressing the challenges posed by high volatility and complex market dynamics. It combines RL's sequential decision-making capability with DL's ability to handle intricate inputs, offering an optimal approach for scenarios requiring both [21][22]. To effectively predict stock prices using time series analysis, it's essential to manage large, complex datasets and make sequential decisions to adapt to the dynamic and uncertain nature of the stock market. Unlike Deep Learning (DL) and Reinforcement Learning (RL), Deep Reinforcement Learning (DRL) is well-suited for such scenarios. DRL's capacity to handle intricate datasets and make sequential decisions in dynamic environments positions it as a powerful tool for predicting stock market trends with high accuracy[23][24][25]. This capability makes DRL highly suitable for addressing the challenges presented by stock market prediction, surpassing conventional DL and RL methods and leading to improved overall forecasting precision

II. RELATED WORK

Therapid advancement of deep learning and reinforcement learning has revolutionized the financial market by automating trading processes [11]. Predicting stock market fluctuations involves dealing with complex nonlinear characteristics, prior decisions, budget constraints, and time limitations. These challenges make decision-making in stock market prediction highly demanding. Traditional statistical models rely on predefined features and assumptions, while machine learning techniques may struggle with handling sequential dependencies in time series data. Although deep learning algorithms can capture intricate patterns, they often require specific domain expertise for feature engineering and may not adapt quickly to rapidly changing market conditions. To address these challenges, we propose a powerful solution: leveraging Deep Reinforcement Learning (DRL). DRL represents a potent combination of reinforcement learning and deep learning techniques. This fusion equips machines with the ability to effectively tackle complex decision-making challenges that were once considered beyond their reach, enabling them to approach real-world problems with human-like intelligence. DeepMind published a groundbreaking algorithm in this field [13]. This approach enables the model to identify intricate patterns that traditional techniques might overlook. DRL's adaptability to changing market conditions ensures that the model remains accurate even in volatile scenarios, enabling timely and informed investment decisions. By incorporating DRL into stock market prediction, our method significantly reduces the reliance on manual feature engineering, streamlining the modeling process and enhancing prediction accuracy. In summary, our proposed approach represents an innovative and effective method for stock market prediction, surpassing traditional statistical models, machine learning techniques, and standalone deep learning algorithms. Through the integration of DRL, we aim to push the boundaries of predictive capabilities in financial markets, facilitating more reliable and efficient decision-making processes.

III. STOCK MARKET PREDICTION ANALYSIS

In the realm of financial prediction, where the goal is to foresee how markets will behave, the tricky task of predicting the unpredictable movements of stock returns has prompted a deep exploration of different ways to tackle it. Specifically, two methods—Deep Learning and Reinforcement Learning—have shown promise, but they also have their limitations when used alone. However, when we merge the strong points of these two methods, a powerful approach emerges: Deep Reinforcement Learning (DRL). As we embark on this journey of exploration, we'll take a close look at each method, pointing out what they do well and where they have shortcomings. This journey ultimately leads us to fully grasp how DRL excels in handling the complex and ever-changing patterns of stock market trends, resulting in more accurate predictions.

A. Deep Learning in Stock Price Prediction

Machine learning is the process of enabling machines to perform tasks without explicit programming. Its fundamental objective is to create algorithms and models capable of making consistently accurate predictions. Within the realm of machine learning, there exists a subset known as Deep Learning (DL). DL endeavors to emulate data consumption processes akin to the human brain. The "Deep" in Deep Learning alludes to the utilization of deep neural networks. While these neural networks may not replicate human cognitive proficiency precisely, they excel at processing vast datasets, making them particularly suitable for handling large-scale datasets. Deep neural networks are constructed with multiple layers, and their effectiveness increases when additional hidden layers are incorporated. The true strength of deep learning lies in its capacity to discern patterns within the input data. The process begins with the provision of a substantial dataset. Deep neural networks then extract pertinent information from this dataset and identify relevant patterns. Subsequently, these deep neural networks endeavor to comprehend and internalize the extracted patterns, which they subsequently utilize for making pertinent predictions. Recurrent Neural Networks (RNNs) exhibit efficacy in handling time series data. Unlike conventional approaches, RNNs adopt a loop-like architecture, facilitating sequential data processing. However, in the context of stock market prediction, relying solely on RNNs proves less effective. Similar to traditional deep learning models, RNNs confront challenges when attempting to decipher long-term stock patterns, rendering them suboptimal for this specific application. In addition to practical implementation considerations, we draw insights from pertinent research papers in the field [29]. These papers propose a novel method for training deep neural networks using reinforcement learning techniques to enhance time series forecasting. By integrating the capabilities of deep learning and reinforcement learning, this approach leads to enhanced accuracy in predicting stock prices [27]. Furthermore, a comprehensive deep learning system was developed for short-term stock market price trend prediction, leveraging deep learning techniques, including RNNs. This system effectively captures temporal dependencies and intricate patterns in historical stock data, resulting in improved prediction performance [25][26]. Another research effort presented a stock price prediction method grounded in deep learning technology. By employing deep neural networks and various deep learning techniques, their approach adeptly captures and analyzes complex patterns within historical stock data, enabling precise predictions of future stock prices [24]. A comparative study conducted by [24] analyzed time series and deep learning algorithms for stock price prediction. This study underscored the strengths of deep learning, particularly RNNs, in capturing intricate patterns and temporal dependencies, highlighting their effectiveness in achieving accurate stock price predictions. Deep learning finds applications in recommendation systems, credit card fraud detection systems, speech recognition, and emerging technologies such as self-driving cars [26][27]. Nonetheless, deep learning may not be the preferred choice when dealing with exceedingly compact datasets. Moreover, the volatile nature of financial markets, influenced by external factors such as geopolitical events and economic news releases, introduces an element of unpredictability that deep learning models may struggle to capture adequately. This unpredictability can lead to occasional discrepancies between model predictions and actual market behavior, further emphasizing the need for robust and adaptive prediction strategies. In addition, the scalability of deep learning approaches can be a consideration in resource-constrained environments, as training large-scale models may require significant computational resources. This aspect necessitates careful consideration when implementing deep learning solutions for stock market prediction, particularly in real-time applications. Furthermore, it is crucial to highlight that while deep learning models can offer impressive accuracy when ample data is available, they may face challenges in low-data or data-scarce scenarios and selecting the right tool for the specific context of stock market prediction.

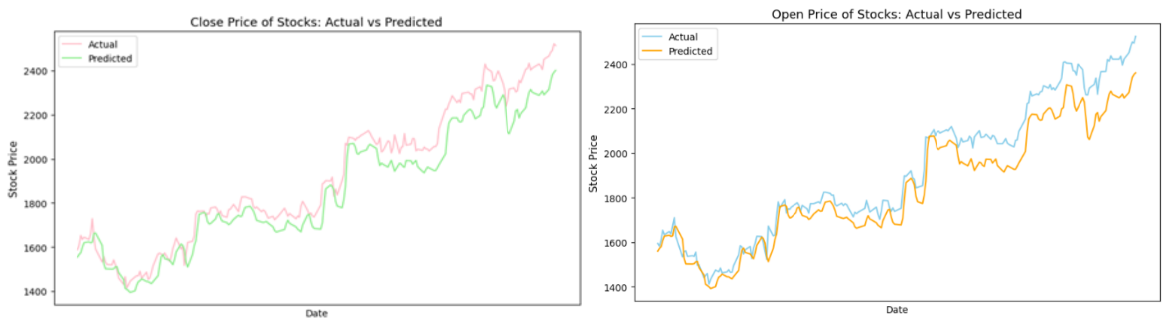

Fig. 1. Stock Price Prediction by Deep Learning

B. Reinforcement Learning in Stock Price Prediction

Reinforcement learning (RL) made its first appearance in the financial market back in 1997 [1]. RL, a branch of machine learning, has found extensive applications in various financial contexts, such as portfolio management for stocks and bonds, as demonstrated by Kanwar in 2019 [2], and Cummings in 2015 for foreign exchange market management [3]. This approach is particularly valuable when an agent interacts with an environment and learns from its experiences, making it well-suited for scenarios where learning through trial and error is essential. In RL, an agent interacts with an environment, making decisions and receiving feedback in the form of rewards or penalties. The primary objective of the agent is to learn a policy that maps states to actions, with the aim of maximizing the expected cumulative reward over time. To achieve this, the agent explores various actions to gather information about the environment while exploiting its current knowledge to make better decisions. The learning process in RL typically revolves around estimating a value function or a Q-function. The value function assesses the expected cumulative reward starting from a specific state and following a particular policy. Conversely, the Q-function estimates the expected cumulative reward for taking a specific action in a given state and following a particular policy thereafter. Among the fundamental RL algorithms, the Q-Learning algorithm is highly regarded and has been widely applied [4]. Reinforcement learning can be categorized into two main types: on-policy learning and off-policy learning, with corresponding algorithms developed for each category [5][6]. Q-learning, a well-known RL technique, seeks to maximize the overall reward [7]. Li et al. [8] applied deep reinforcement learning to stock forecasting, harnessing the power of Q-learning. Both Q-learning and SARSA methods have been developed to align with the adaptive market hypothesis, which considers the relatively slow adoption of new information by the market compared to its arrival [9]. Deep Q-learning has been employed in contexts such as the foreign exchange market in comparison to buy and hold strategies and expert traders [10], as well as in stock market indices [11].

Exploration and exploitation are pivotal concepts in RL. Exploration involves the agent's experimentation with different actions to gain a better understanding of the environment, while exploitation entails utilizing the agent's current knowledge to make decisions likely to yield high rewards. Striking the right balance between exploration and exploitation is crucial for success in RL. Furthermore, applying RL to time series data involves training agents to optimize decisions over time. This encompasses modeling the problem as a Markov Decision Process (MDP), estimating value functions, employing Temporal Difference (TD) learning, utilizing Recurrent Neural Networks (RNNs) to capture temporal dependencies, and implementing experience replay. Challenges in this domain include handling sequential correlations, variable-length time series, and non-stationary environments. Various techniques have been proposed to address these challenges and enhance the performance of time series RL algorithms. Deep reinforcement learning takes RL a step further by incorporating deep neural networks to approximate value functions or policies. These networks empower the agent to learn from high-dimensional input, such as images or sensor data and can handle more complex tasks, including stock price prediction and forecasting, compared to traditional RL methods.

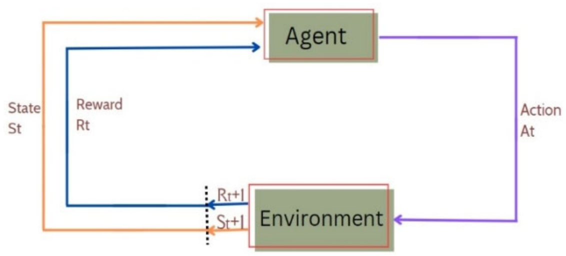

Fig. 2. Reinforcement Learning Architecture

C. Deep Reinforcement Learning in Stock Price Prediction

Deep Deep Reinforcement Learning (DRL) stands as a captivating subfield within Artificial Intelligence (AI), where it seamlessly integrates deep learning techniques. Deep learning, a subset of Machine Learning, centers around neural networks to make supervised predictions from input data, encoding valuable data representations [30]. DRL also incorporates Reinforcement Learning, an interaction-based learning approach where an agent learns through trial and error within an environment [31]. This trial-and-error process is foundational to how reinforcement learning agents acquire knowledge [31]. DRL empowers deep neural networks to comprehend their surroundings effectively, enabling autonomous agents to make informed decisions. These agents can tackle complex tasks through neural networks in high-dimensional input spaces. In the context of time series data, DRL leverages a Deep Q-learning algorithm to train agents to learn optimal strategies. This algorithm estimates the best action-value function for a given state, guiding the agent's actions to maximize rewards over time [31][32]. This approach holds promise for various time series applications, such as Stock Price Prediction. DRL combines the strengths of sequential decision-making from reinforcement learning with deep learning's pattern modeling capability, allowing it to capture both short-term fluctuations and long-term trends in stock prices. Typically, deep reinforcement learning models for stock price prediction employ Recurrent Neural Networks (RNNs) or variants like Long Short-Term Memory (LSTM) networks to handle sequential data. These networks effectively capture temporal dependencies and patterns in stock price movements. Reinforcement learning algorithms, such as Q-learning or policy gradients, are then integrated to optimize the agent's trading strategy. Numerous approaches in deep reinforcement learning have been proposed for stock price prediction. For example, Liang et al. [38] and Spooner et al. [39] explore adversarial deep reinforcement learning and market making via reinforcement learning, respectively. Xiong et al. [34] introduce a practical deep reinforcement learning approach for stock trading, while Yu et al. [35] propose a model-based deep reinforcement learning approach for dynamic portfolio optimization. Carta et al. [36] suggest a multi-layer and multi-ensemble stock trader utilizing deep learning and deep reinforcement learning. Aboussalah and Lee [37] introduce an innovative portfolio optimization method using continuous control strategies combined with stacked deep dynamic recurrent reinforcement learning.

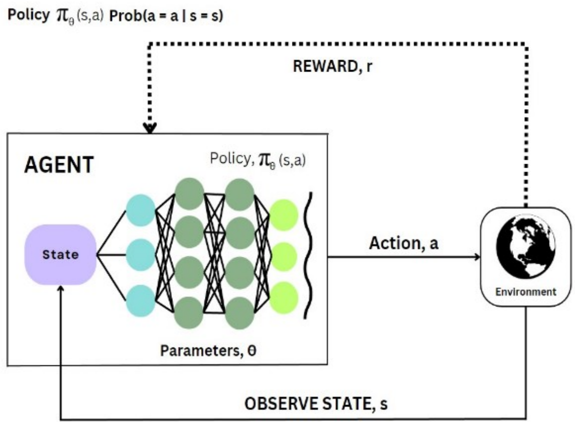

Fig. 3. Deep Reinforcement Learning Model Architecture

In our approach, the central focus lies on leveraging Deep Q-Learning for precise stock price prediction. This algorithm marries the strengths of deep learning and reinforcement learning, functioning as a model-free technique where an agent autonomously learns optimal actions across diverse scenarios. Within the realm of reinforcement learning, the fundamental aim is to cultivate an optimal policy that maximizes cumulative rewards. To accomplish this, we employ Q-learning, a method that employs Q-values to represent the expected future rewards associated with specific actions within particular states. Deep Q-

Learning introduces a pivotal component: the Q-network. This deep neural network takes the current state of the environment as input and yields corresponding Q-values for all available actions. Throughout the training process, the Q-network relentlessly strives to minimize the discrepancy between its predicted Q-values and the target Q-values, achieved through the application of the Q-learning update rule. This iterative refinement process ultimately equips the agent with the capability to make strategic decisions that maximize expected returns. Empirical results substantiate the efficacy of Deep Reinforcement Learning, particularly when employing the Deep Q-Network (DQN) approach for stock price prediction. This approach consistently outperforms traditional Deep Neural Networks (DNNs) and standalone Reinforcement Learning techniques [40][41]. The symbiosis of deep learning and reinforcement learning methodologies within the DRL framework significantly empowers agents to excel in tackling intricate tasks such as stock price prediction. This integration enhances the decision-making prowess of agents, enabling them to thrive in dynamic and challenging environments [38][40][41][42]. The holistic approach of combining deep learning's pattern recognition capabilities with reinforcement learning's sequential decision-making expertise creates a robust foundation for addressing complex financial market prediction tasks, showcasing the potential of DRL as a transformative tool in this domain.

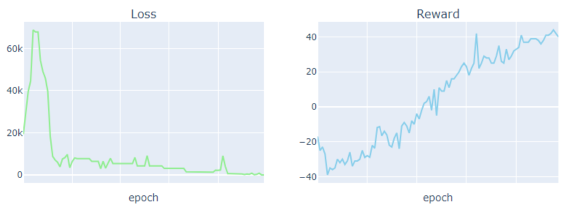

Fig. 4. Performance Trends: Loss and Reward (DRL)

IV. RESULTS AND DISCUSSIONS

The experimental evaluations of our study provide solid proof of the effectiveness of our integrated approach in predicting stock market movements. By combining traditional time series techniques with Deep Reinforcement Learning (DRL), we have surpassed both conventional methods and standalone deep learning models. The integration of DRL has streamlined the prediction system, reducing the need for manual feature engineering and thus minimizing potential human errors and biases. Reducing maintenance costs and improving overall prediction accuracy, the DRL model's remarkable adaptability to the ever-changing dynamics of the stock market is a significant advantage. Its decision-making strategies can swiftly adjust as market conditions fluctuate. This continuous learning from the latest market data enables it to capture complex patterns and dependencies, ultimately leading to more precise and up-to-date forecasts. Combining traditional time series techniques with DRL creates a powerful, cost-effective solution for accurate stock market predictions. This integration's reduction in manual feature engineering and the DRL model's adaptability give it a competitive edge in the dynamic and highly competitive world of stock market prediction. The DRL model's exceptional ability to react to changing market conditions opens up new possibilities for predictive modeling in the financial industry and other areas beyond stock market forecasting.

V. CONCLUSION

In conclusion, our research paper delved into the application of reinforcement learning in stock prediction, to navigate the uncertainties of financial markets, achieving a balance between exploration and exploitation, and enhancing prediction accuracy for maximizing profits. Through our investigation, we found that reinforcement learning, especiallyDeep Reinforcement Learning (DRL), effectively adapts to the dynamic nature of stock markets and reduces the dependency on manual feature engineering, leading to a more efficient and accurate prediction system. Looking ahead, we propose several avenues for future work to further enhance stock prediction accuracy: a) Analyzing and combining diverse exploration methods to devise novel approaches for effective exploration in the stock market context. b) Training a curious agent in multiple environments without rewards to explore the transferability of learning to target environments with rewards. c) Exploring global exploration strategies that entail coordinated decisions over extended time horizons to optimize decision-making in complex market conditions. Our research primarily focused on enhancing stock prediction accuracy using DRL within the context of time series analysis. The developed DRL model showcased promising results by significantly reducing manual feature engineering, thereby saving time and resources while enhancing prediction accuracy. Using these insights, future research can enhance DRL's stock market forecasting and broader financial applications.

REFERENCES

[1] Moody, J., Wu, L. Liao, Y., Saffell, M. Performance functions and reinforcement learning for trading systems and portfolios. J. Forecast. 1998, 17, 441–470.

[2] Kanwar, N. Deep Reinforcement Learning-Based Portfolio Management; The University of Texas at Arlington: Arlington, TX, USA, 2019.

[3] Cumming, J. An Investigation into the Use of Reinforcement Learning Techniques within the Algorithmic Trading Domain; Imperial College London: London, UK, 2015.

[4] Dasgupta, Agnibh. Deep Q Learning Applied to Stock Trading. Diss. Utah State University, 2020.

[5] Pendharkar, P.C.; Cusatis, P. Trading financial indices with reinforcement learning agents. Expert Syst. Appl. 2018, 103, 1–13. [CrossRef]

[6] D'Eramo, C.; Restelli, M.; Nuara, A. Estimating the Maximum Expected Value through Gaussian Approximation. Int. Conf. Mach. Learn. 2016, 48, 1032–1040.

[7] C. Jin, Z. Allen-Zhu, S. Bubeck, M.I. Jordan, Is Q-learning provably efficient? in Advances in 174 neural information processing systems (2018), pp. 4863–4873

[8] Y. Li, P. Ni, V. Chang, Application of deep reinforcement learning in stock trading strategies and stock forecasting. Computing 1–18 (2019).

[9] Marco Corazza, A.S. Q-Learning and SARSA: A comparison between two intelligent stochastic control approaches for financial trading. Univ. Ca' Foscari Venice Dept. Econ. Res. Pap. 2015, 15, 1–23.

[10] Sornmayura, S. Robust forex trading with deep q network (dqn). Assumpt. Bus. Adm. Coll. 2019, 39, 15–33.

[11] Mosavi A.et al. A comprehensive review of deep reinforcement learning methods and applications in economics. Mathematics, 2020;8.10: 1640.

[12] Lee, J.W.; Park, J.; O, J.;Lee, J.; Hong, E. A Multiagent Approach to Q-Learning for Daily Stock Trading. IEEE Trans. Syst. ManCybern. -Part A Syst. Hum. 2007, 37, 864–877.

[13] Mnih V, Kavukcuoglu K, Silver D, Rusu AA, Veness J et al. Human-level control through deep reinforcement learning. Nature 2015; 518 (7540): 529-533.

[14] C.K.S. Leung, R.K. MacKinnon, Y. Wang, A machine learning approach for stock price prediction, in Proceedings of the 18th International Database Engineering & Applications Symposium (2014), pp. 274–277

[15] Hryshko, A.; Downs, T. System for foreign exchange trading using genetic algorithms and reinforcement learning. Int. J. Syst. Sci. 2004, 35, 763–774.

[16] Dang, T. H., et al. (2019). Reinforcement Learning for Financial Trading: A Review. Journal of Financial Data Science, 1(2), 68-83.

[17] Chen, Y., et al. (2021). Reinforcement Learning for Stock Trading: A Comprehensive Review. arXiv preprint arXiv:2106.13976.

[18] Moerland, T. M., et al. (2019). Reinforcement Learning for Trading under Limit Order Book Dynamics. Proceedings of the 28th International Joint Conference on Artificial Intelligence (IJCAI-19).

[19] Bengio, Y. Learning deep architectures for AI. Foundations and Trends in Machine Learning 2, 1–127 (2009).

[20] Lin, W., et al. (2021). Financial Time Series Prediction using Deep Reinforcement Learning with Technical Analysis and Fundamentals. Expert Systems with Applications, 177, 114947.

[21] Zhang, X., et al. (2020). Interpretability of Deep Reinforcement Learning Stock Trading Strategies. Expert Systems with Applications, 143, 113097.

[22] I. Parmar, N. Agarwal, S. Saxena, R. Arora, S. Gupta, H. Dhiman, L. Chouhan, Stock market prediction using machine learning, in 2018 First International Conference on Secure Cyber Computing and Communication (ICSCCC) (IEEE, 2018), pp. 574–576.

[23] P. Rani, J. Shokeen, D. Mullick, Recommendations using modified k-means clustering and voting theory. Int. J. Comput. Sci. Mobile Compute.

[24] J. Shokeen, C. Rana, Social recommender systems: techniques, domains, metrics, datasets and future scope. J. Intell. Inform. Syst. 54, 633–667 (2019). https://doi.org/10.1007/s10844-019-00578-5.

[25] J. Shokeen, C. Rana, A study on features of social recommender systems. Artif. Intell. Rev.53(2), 965–988 (2020).

[26] J. Shokeen, C. Rana, P. Rani, A trust-based approach to extract social relationships for recommendation, in Data Analytics and Management (Springer, 2020), pp. 51–58.

[27] Takashi Kuremoto, Takaomi Hirata, Masanao Obayashi, Shingo Mabu, and Kunikazu Kobayashi. "Training Deep Neural Networks with Reinforcement Learning for Time Series Forecasting.".

[28] Jingyi Shen and M. Omair Shafq. "Short-term Stock Market Price Trend Prediction using a Comprehensive Deep Learning System.".

[29] Xuan Ji, Jiachen Wang, and Zhijun Yan. "A Stock Price Prediction Method Based on Deep Learning Technology. IEEE Trans. Syst. ManCybern. -Part A Syst. Hum. 2007, 37, 864–877.

[30] Santosh Ambaprasad Sivapurapu. "Comparative Study of Time Series and Deep Learning Algorithms for Stock Price Prediction."

[31] Will Serrano. Deep reinforcement learning with the random neural network. Engineering Applications of Artificial Intelligence, 110:104751, 2022.

[32] Anil Berk Altuner and Zeynep Hilal Kilimci. A novel deep reinforcement learning based stock direction prediction using knowledge graph and community aware sentiments, 2021.

[33] Xiang Gao. Deep reinforcement learning for time series: playing idealized trading games, 2018.

[34] Xiong, Zhuoran, Xiao-Yang Liu, Shan Zhong, Hongyang Yang, and Anwar Walid. Practical deep reinforcement learning approach for stock trading. arXiv preprint arXiv:1811.07522, 2018.

[35] Yu, Pengqian, Joon Sern Lee, Ilya Kulyatin, Zekun Shi, and Sakyasingha Dasgupta. Model-based deep reinforcement learning for dynamic portfolio optimization. arXiv preprint arXiv:1901.08740, 2019.

[36] Carta, Salvatore, Andrea Corriga, Anselmo Ferreira, Alessandro Sebastian Podda, and Diego Reforgiato Recupero. A multi-layer and multi-ensemble stock trader using deep learning and deep reinforcement learning. Applied Intelligence, 2021;51(2):889-905.

[37] Aboussalah, Amine Mohamed, and Chi-Guhn Lee. Continuous control with stacked deep dynamic recurrent reinforcement learning for portfolio optimization. Expert Systems with Applications, 2020;140: 112891.

[38] Liang, Zhipeng, Hao Chen, Junhao Zhu, Kangkang Jiang, and Yanran Li. Adversarial deep reinforcement learning in portfolio management. arXiv preprint arXiv:1808.09940 (2018).

[39] Spooner T., Fearnley J. Savani R., and Koukorinis A. Market making via reinforcement learning. arXiv preprint arXiv:1804.04216, 2018.

[40] Rangappa Jyothi and Gorappa Ningappa Krishnamurthy. Deep-reinforcement learning-based architecture for multi-objective optimization of stock prediction. European Journal of Electrical Engineering and Computer Science, 6(4):9–16, 2022.

[41] Li, Y., Ni, P., & Chang, V. (2019). Application of deep reinforcement learning in stock trading strategies and stock forecasting. Computing, 1-18.

[42] R.S. Sutton, A.G. Barto, Reinforcement Learning: An Introduction (MIT Press, 2018) Deep Learning: A Comprehensive Overview on Techniques, Taxonomy, Applications, and Research Directions. (2021, August 18).

[43] V.K.S. Reddy, Stock market prediction using machine learning. Int. Res. J. Eng. Technol. 5(10)(2018.

★★★